THE 25-YEAR-OLD INTEL PATENT THAT KILLED YOUR DISASSEMBLERS ft. colby57

Hello everyone, it's been a while since my last post, and today we're going to fix that.

Our story began with a message from my friend-reverser colby57: *"The first time I saw this, I thought I was about to crash, but everything was valid"*. The attachment was a shell sample. Nothing special at first glance, but in the middle of the code, the disassembler simply broke off, showing "garbage" bytes further on. The processor, however, executed this code without any issues.

This "garbage" turned out to be two opcodes: `0F 1A` and `0F 1B`. There are actually more of them, but only two will be mentioned.

And this isn't a bug in a specific sample. It's a systemic failure of the entire reverse engineering industry, laid down by an Intel patent almost 30 years ago.

Digging into the past: where did these ghosts come from?

The roots of our problem go back to 1997, to one of Intel's most elegant patents, created at a time when the world was preparing for the Y2K problem. Back then, Intel faced an eternal headache - backward compatibility. Engineers added bunches of cool instructions to new processors, but developers wouldn't touch them, fearing their programs wouldn't run on computers released a couple of years earlier.

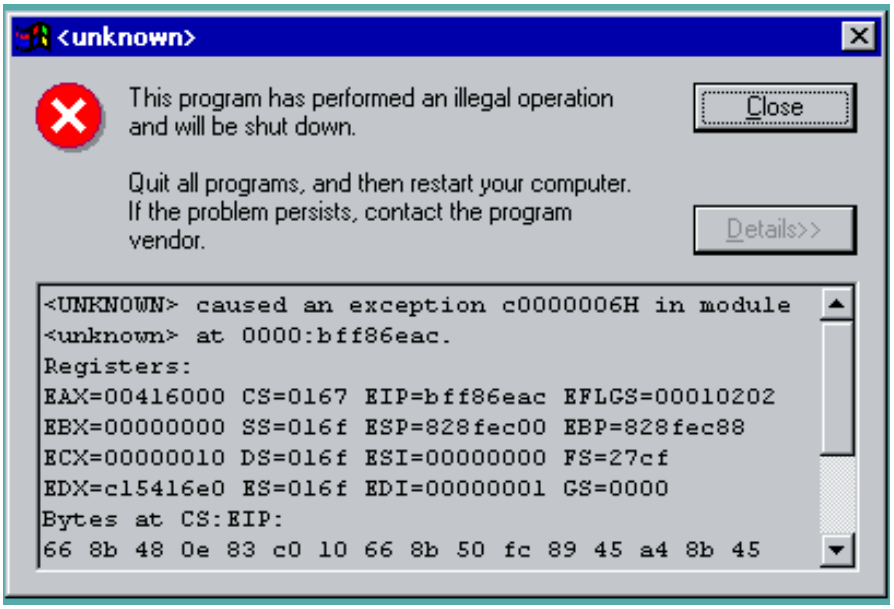

The solution proposed in patent **US5,701,442** was brilliantly simple. "What if we reserve an entire range of opcodes - from `0x0F18` to `0x0F1F` - as just placeholders?" In the patent, they were called **Hintable NOPs**.

On processors of that time, they did absolutely nothing. But at any moment, any placeholder could "come alive" and turn into a new, useful instruction. On older computers, such an instruction would execute as a harmless NOP, on newer ones - as a powerful tool.

And so it happened. `0x0F18` turned into the **PREFETCH** family of instructions, and `0x0F1E` after twenty years became the heart of CET protection technology in the form of **ENDBR64**. Most opcodes from this range were eventually either utilized or, at the very least, correctly recognized by disassemblers.

Except for two. Opcodes `0F 1A` and `0F 1B` became real ghosts. They remained in the shadows, forgotten by everyone except the silicon itself and the pages of that very patent.

Live test: breaking everything

Enough theory. My good friend wrote a simple PoC binary containing these instructions, and we decided to run it through our pack of standard tools. Let's look at this together.

GitHub repository with PoC

First patient: x64dbg

We load our binary into the debugger. Here they are, our "garbage" bytes. x64dbg honestly shows us `???`, unable to recognize the instruction.

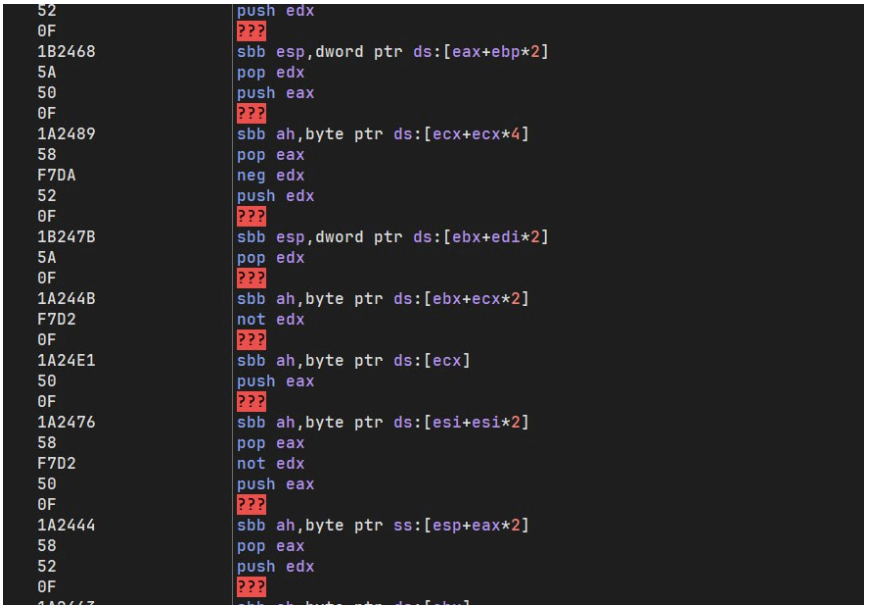

The EIP arrow points to the line with bytes 0F 1A C0. In the disassembler column on this line are three question marks.

But the most interesting thing is the processor's behavior. We step (F7), and... EIP calmly jumps over this "unknown" instruction and continues. No crash, no error. For the CPU, it's valid, executable code. For the debugger - some nonsense

Second patient: IDA Pro

And what will our beloved IDA say? We open the binary and see a depressing picture.

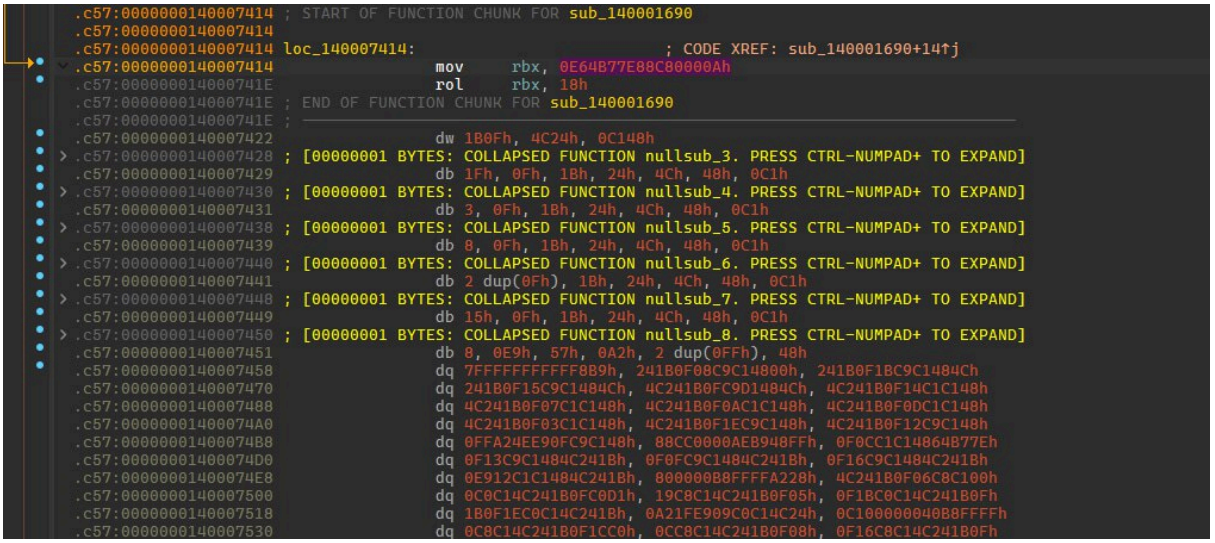

IDA doesn't just not understand the instruction. It decides that **the executable code ends here** and data begins. The function analysis breaks off. All subsequent code simply ceases to exist for the static analyzer. This is very bad, but it's business as usual for it.

Third patient: Binary Ninja & Ghidra

No miracle happened here either. Binary Ninja marks the instruction as `??`. Ghidra shows `undefined`. The same story - the industry's leading tools see valid code as an error

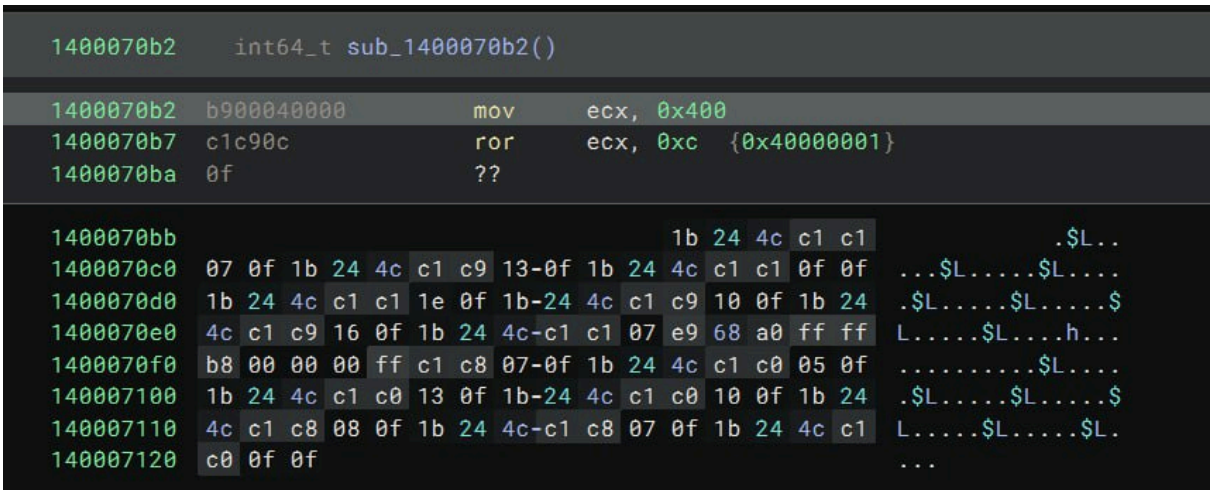

How can we use this?

So what do these two opcodes give an attacker in practice?

The most obvious - **breaking analysis**. As we saw in IDA, one such NOP - and the function analysis breaks off. All subsequent code simply ceases to exist for the disassembler.

A higher-level example: **creating fake execution graphs**. After a `jmp`, a combination of `0F 1B E8...` is placed, which the disassembler will mistakenly take as a `CALL`, sending the analyst down a false trail that leads nowhere.

And this is just the tip of the iceberg; thousands of different applications can be devised.

An industry problem. A human problem.

But all these tricks are just consequences. The real problem is that a fundamental, documented part of the x86 architecture remained invisible to our most advanced tools for **decades**.

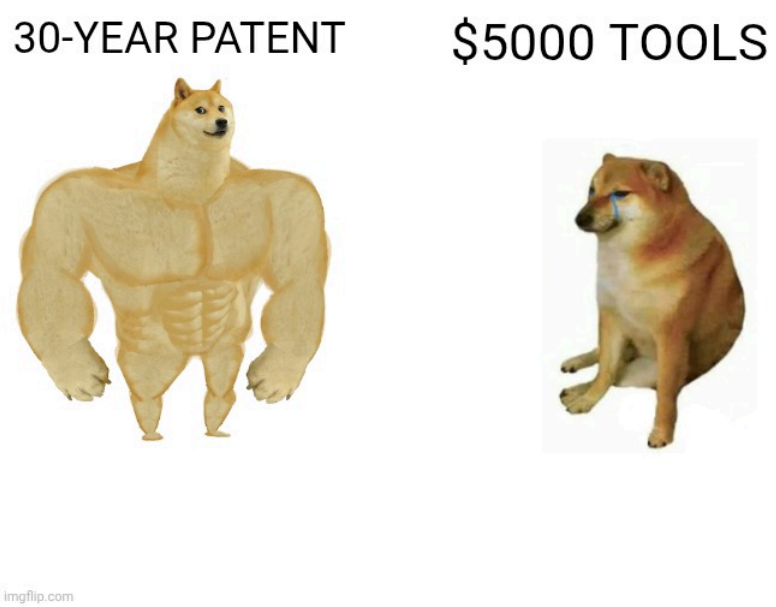

Think about it. This isn't some exotic feature from secret manuals. Not an undocumented feature of a new processor. These are basic instructions described in a patent almost 30 years ago. And all this time they just... existed. Right under our noses.

The sample that started it all revealed not a vulnerability in the processor, but a huge problem in our approach.

That's all for now. Thanks for reading.